

Home

Assignment Regression

localhost:8888/notebooks/Assignment%20Regression.ipynb?

Jupyter Assignment Regression Last Checkpoint: yesterday

File Edit View Run Kernel Settings Help

Trusted

Markdown

JupyterLab Python 3 (ipykernel)

### REGRESSION

[ ]: 

```
# The objective of this assignment is to evaluate your understanding of regression techniques in supervised learning by
# applying them to a real world dataset
```

[ ]: 

```
# The dataset california housing available in the sklearn library contains information about various features of houses in
# California and their respective median prices
```

[6]: 

```
import pandas as pd
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

### Loading and Preprocessing

[8]: 

```
# Load the California Housing Dataset
Cal_Housing = fetch_california_housing()
```

[10]: 

```
# Convert the dataset into a pandas Dataframe
df = pd.DataFrame(Cal_Housing.data, columns = Cal_Housing.feature_names)
df['Target']=Cal_Housing.target
```

[19]: df

[19]:

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	Target
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23	4.526
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22	3.585
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24	3.521
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	3.413

29°C Sunny

Search

ENG IN

11:31 05-01-2025

Home

Assignment Regression

localhost:8888/notebooks/Assignment%20Regression.ipynb?

Assignment Regression Last Checkpoint: yesterday

Trusted

File Edit View Run Kernel Settings Help

JupyterLab Python 3 (ipykernel)

0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23	4.526
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22	3.585
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24	3.521
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	3.413
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25	3.422
...	...	...	...	...	...	...	...	...	...
20635	1.5603	25.0	5.045455	1.133333	845.0	2.560606	39.48	-121.09	0.781
20636	2.5568	18.0	6.114035	1.315789	356.0	3.122807	39.49	-121.21	0.771
20637	1.7000	17.0	5.205543	1.120092	1007.0	2.325635	39.43	-121.22	0.923
20638	1.8672	18.0	5.329513	1.171920	741.0	2.123209	39.43	-121.32	0.847
20639	2.3886	16.0	5.254717	1.162264	1387.0	2.616981	39.37	-121.24	0.894

20640 rows x 9 columns

[21]: df.head()

[21]:

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	Target
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23	4.526
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22	3.585
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24	3.521
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	3.413
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25	3.422

29°C Sunny

Search

ENG IN

11:31 05-01-2025



3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	3.413
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25	3.422

[27]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  -
0   MedInc      20640 non-null  float64
1   HouseAge    20640 non-null  float64
2   AveRooms    20640 non-null  float64
3   AveBedrms   20640 non-null  float64
4   Population  20640 non-null  float64
5   AveOccup    20640 non-null  float64
6   Latitude    20640 non-null  float64
7   Longitude   20640 non-null  float64
8   Target      20640 non-null  float64
dtypes: float64(9)
memory usage: 1.4 MB
```

[12]: # Handling missing values  
print('Missing Values Count:', df.isnull().sum())

```
Missing Values Count: MedInc      0
HouseAge      0
AveRooms      0
AveBedrms     0
Population    0
AveOccup      0
Latitude      0
Longitude     0
Target        0
dtype: int64
```



Home

Assignment Regression

localhost:8888/notebooks/Assignment%20Regression.ipynb?

Jupyter Assignment Regression Last Checkpoint: yesterday

Trusted

File Edit View Run Kernel Settings Help

JupyterLab Python 3 (ipykernel)

[14]: # Perform Feature Scaling using Standard Scaler

scaler = StandardScaler()

scaled\_data = scaler.fit\_transform(df[Cal\_Housing.feature\_names])

scaled\_data

[14]:

array([[ 2.34476576, 0.98214266, 0.62855945, ..., -0.04959654,

1.05254828, -1.32783522],

[ 2.33223796, -0.60701891, 0.32704136, ..., -0.09251223,

1.04318455, -1.32284391],

[ 1.7826994 , 1.85618152, 1.15562047, ..., -0.02584253,

1.03850269, -1.33282653],

...,

[-1.14259331, -0.92485123, -0.09031802, ..., -0.0717345 ,

1.77823747, -0.8237132 ],

[-1.05458292, -0.84539315, -0.04021111, ..., -0.09122515,

1.77823747, -0.87362627],

[-0.78012947, -1.00430931, -0.07044252, ..., -0.04368215,

1.75014627, -0.83369581]])

[16]: #Split the dataset into training and testing sets

X\_train , X\_test , Y\_train , Y\_test = train\_test\_split(df[Cal\_Housing.feature\_names],df['Target'],test\_size=0.2, random\_state = 42)

[18]: print('Preprocessed DataFrame:')

print(df.head())

Preprocessed DataFrame:

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	\
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	

	Longitude	Target
0	-122.23	4.526
1	-122.22	3.585
2	-122.24	3.521

29°C Sunny

Search

ENG IN

11:32 05-01-2025

4. Feature Scaling by using StandardScaler , This is necessary because some features have different units or scales which affect the performance of machine learning models.By scaling these features to have zero mean and unit variance ,we can improve the stability and accuracy of the model.



## Regression Algorithm Implementation

### Linear Regression

```
[22]: from sklearn.linear_model import LinearRegression
      model = LinearRegression()
      model.fit(X_train , Y_train)
      y_predict = model.predict(X_test)
```

```
[26]: y_predict
```

```
[26]: array([0.71912284, 1.76401657, 2.70965883, ..., 4.46877017, 1.18751119,
          2.00940251])
```

Linear Regression is a linear model that predicts continuous output variable based on one or more input features. It works by learning the coefficients of a linear equation that minimizes the difference between predicted and actual values. Linear Regression is suitable for this dataset because it's a simple and interpretable model that can capture linear relationships between features and the target variables. It may not perform well if the relationships are linear.

### Model Evaluation and Comparison

```
[45]: from sklearn.metrics import mean_squared_error , mean_absolute_error , r2_score
```

```
[47]: #Performance of Linear Regression algorithm using metrics
      MSE = mean_squared_error (Y_test,y_predict)
      MAE = mean_absolute_error (Y_test,y_predict)
      r2 = r2_score (Y_test,y_predict)
```

```
print("Linear Regression Performance Metrics:")
print(f"Mean Squared Error (MSE): {MSE:.2f}")
print(f"Mean Absolute Error (MAE) : {MAE:.2f}")
print(f"R-Squared Score (r2) : {r2:.2f}")
```

Linear Regression Performance Metrics:





```
[22]: from sklearn.linear_model import LinearRegression
      model = LinearRegression()
      model.fit(X_train, Y_train)
      y_predict = model.predict(X_test)
```

```
[26]: y_predict
```

```
[26]: array([0.71912284, 1.76401657, 2.70965883, ..., 4.46877017, 1.18751119,
          2.00940251])
```

Linear Regression is a linear model that predicts continuous output variable based on one or more input features. It works by learning the coefficients of a linear equation that minimizes the difference between predicted and actual values. Linear Regression is suitable for this dataset because it's a simple and interpretable model that can capture linear relationships between features and the target variables. It may not perform well if the relationships are linear.

### Model Evaluation and Comparison

```
[45]: from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

```
[47]: #Performance of Linear Regression algorithm using metrics
```

```
MSE = mean_squared_error(Y_test, y_predict)
MAE = mean_absolute_error(Y_test, y_predict)
r2 = r2_score(Y_test, y_predict)
```

```
print("Linear Regression Performance Metrics:")
print(f"Mean Squared Error (MSE): {MSE:.2f}")
print(f"Mean Absolute Error (MAE) : {MAE:.2f}")
print(f"R-Squared Score (r2) : {r2:.2f}")
```

```
Linear Regression Performance Metrics:
Mean Squared Error (MSE): 0.56
Mean Absolute Error (MAE) : 0.53
R-Squared Score (r2) : 0.58
```